

# Final Project: Build a Transformer-Based Fake News Detection Chatbot

In this project, you are required to design and deploy an intelligent chatbot that detects fake news using a transformer-based model. The primary aim is to apply your skills in data preprocessing, model integration with a transformer model, and deployment to create a modern AI-driven conversational system for fake news detection.

## Objective

The chatbot should process user-input news text and predict whether it is fake or true based on provided datasets ([true.csv](#) and [fake.csv](#)). The system will serve as a practical application of transformer models to classify news and deliver results conversationally.

## Data

CSV files: [true.csv](#) and [fake.csv](#), containing true and fake news articles, respectively. Unlike the previous assignment, which relied on keyword-based labeling without contextual information, this project requires preprocessing the data to create meaningful embeddings using a transformer model for classification.

## Data Preprocessing

- Combine the [true.csv](#) and [fake.csv](#) datasets into a single dataset with labels (e.g., 1 for true, 0 for fake).
- Clean the text data by removing special characters, punctuation, and stopwords, and perform tokenization suitable for the selected transformer model.
- Ensure the dataset is balanced or address any imbalances during model training.
- Split the dataset into training, validation, and test sets (e.g., 80-10-10 split).
- For less model run time, use only top 100 rows of data from each csv file.(this is not mandatory but recommended)

## Model Selection and Integration

You could use an open-source transformer model from Hugging Face's [transformers](#) library. For this project, we recommend using pre-trained models such as BERT ([bert-base-uncased](#)) or DistilBERT ([distilbert-base-uncased](#)), as they are well-suited for text classification tasks like fake news detection and are computationally efficient for training

on Kaggle's T4 free GPU. BERT provides robust performance, while DistilBERT offers a lighter, faster alternative with comparable accuracy. You may use the model in an inference-only manner after fine-tuning by embedding the input text and generating predictions dynamically. We are open to innovative approaches, such as experimenting with model configurations or optimization techniques.

## Steps for Model Integration

1. Load the pre-trained BERT or DistilBERT model and tokenizer from Hugging Face.
2. Preprocess the text data to create input embeddings compatible with the selected model.
3. Fine-tune the model on the labeled dataset for binary classification (fake vs. true news).
4. Save the trained model for use in the chatbot.

## Chatbot Development

Develop a chatbot interface capable of processing a single user-input news article or text snippet and generating a prediction (fake or true) with a confidence score. The chatbot should use the fine-tuned BERT or DistilBERT model to classify the input and return a conversational response (e.g., "This news is likely fake with a confidence of 85%").

## Deployment with Streamlit

Deploy the chatbot using Streamlit. The interface should include:

- A text input box for users to paste or type a news article/snippet.
- An area to display the model's prediction (fake or true) and confidence score.

## Video Demonstration

Record a short video (under 1 minute) demonstrating your chatbot in action. The video should clearly show:

- The Streamlit interface.
- A few example inputs (e.g., a fake news snippet and a true news snippet).
- The corresponding predictions and confidence scores generated by the chatbot.

## Deliverables and Evaluation

Your submission must be pushed to your GitHub repository and include:

- Source code, including data preprocessing, model training, and chatbot implementation.
- The processed dataset (or scripts to preprocess the provided CSVs).

- Upload your streamlit deploy link to your github repository.
- A README file explaining how to run the project, including dependencies and instructions for running the Streamlit app locally or accessing a public/shareable link.
- A video file (max 1 minute) demonstrating the chatbot functionality.

Evaluation will be based on:

- Effectiveness of data preprocessing and model training.
- Accuracy and relevance of the chatbot's predictions.
- Quality and usability of the Streamlit interface.
- Completeness of deployment and adherence to submission guidelines.

## **Deadline**

The deadline to submit the project is 6th July 2025.